

Deep Learning

MSDS 631

Convolutional Neural Networks

Michael Ruddy

Questions?

- From last lecture?

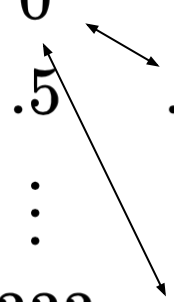
Overview

- What/Why is a Convolution?
- CNN-specific hyperparameters
- Basic CNN history/set-up

Why are images special?

- Images are deceptively hard
- Images are big
- Geometry matters!
 - Pixels near each other interact in different ways to create features than pixels far away
 - This is free data that we lose if we simply consider an image as a data vector

Different
relationships

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$


The Convolution

- Fancy **linear** operation useful for spatial data

The Convolution

- Fancy **linear** operation useful for spatial data

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (.5 \times 0) + (0 \times 0) + (.25 \times 2) \\ = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

=

$$\begin{bmatrix} 1.5 & \dots & \dots \\ \vdots & \ddots & \vdots \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(.5 \times 1) + (1 \times 0) + (.25 \times 0) + (.5 \times 2) = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & 1.5 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (0 \times 0) + (.5 \times 0) + (1 \times 2) = 3$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & 1.5 & \rightarrow & 3 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(0 \times 1) + (.25 \times 0) + (1 \times 0) + (.25 \times 2) = .5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*



$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & ? & ? \\ ? & ? & ? \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & .25 & 2.5 \\ 1 & 2.25 & 2 \end{bmatrix}$$

The Convolution


- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Result}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

- Only four parameters!
 - If input is dimension 16 and output is dimension 9, how many for FC?

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Output Image}$$

Unknown Parameters

“2x2 Filter”

Why Convolution?


- Only four parameters!
- Translational Equivariance
 - If I shift my image, I shift the output!

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Image}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?


- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

Intuition: Edge
Detection

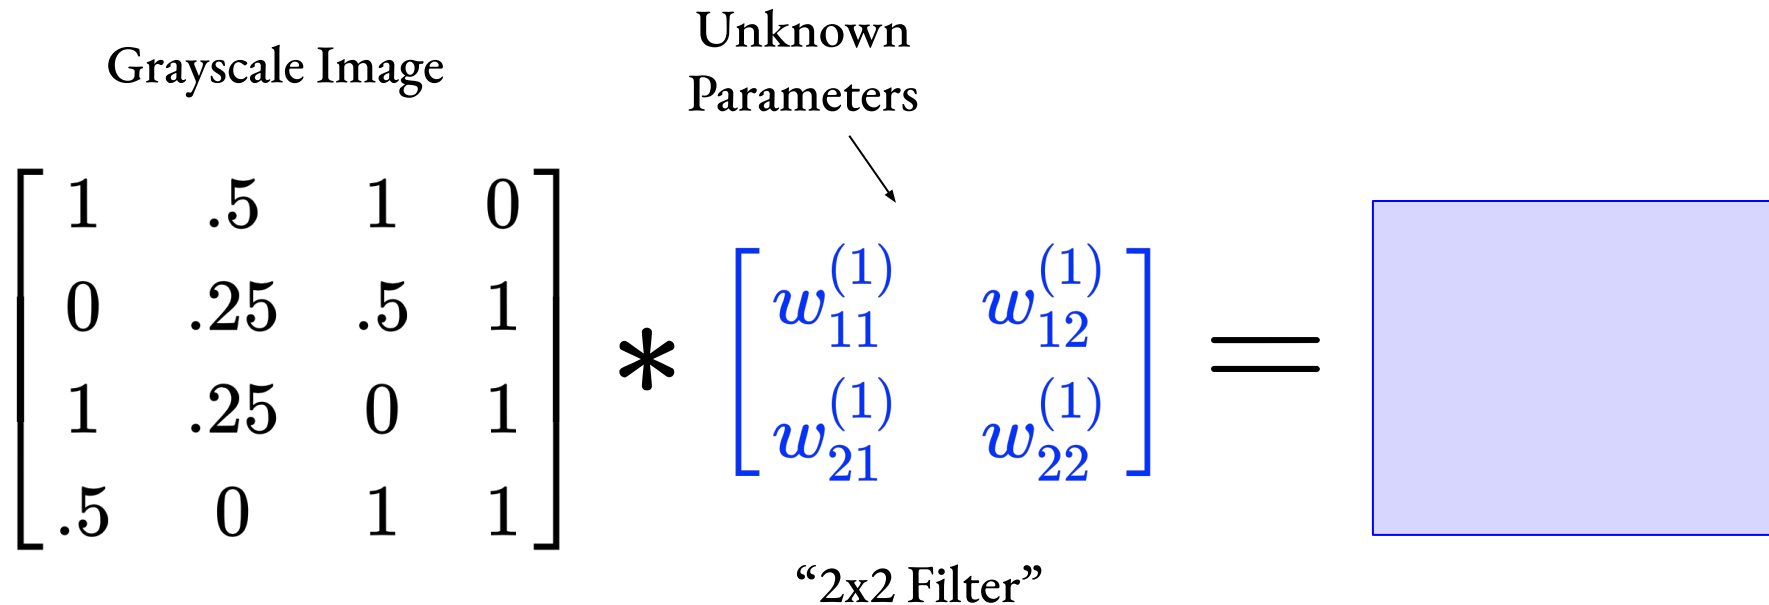
- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

“2x2 Filter”



The Convolution

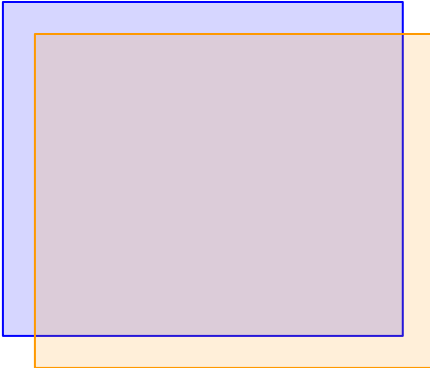
- In a Conv. layer we apply many filter to get many features

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} \end{bmatrix} =$$

More Parameters

“2x2 Filter”



The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

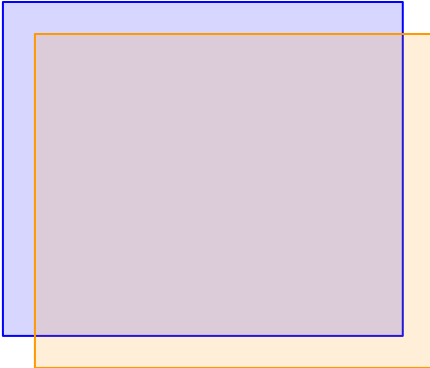
Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} \end{bmatrix} =$$

More Parameters


“2x2 Filter”

“2 Channels”





Layer 3



Layer 2



Layer 1

*Convolutional Deep Belief Networks
for Scalable Unsupervised Learning
of Hierarchical Representations, Lee
H., Grosse R., Ranganath R., Ng A.*

The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

RGB Image

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} .2 \\ .2 \\ 1 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .5 \\ .25 \\ .25 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .15 \\ .25 \\ 1 \\ .5 \\ .75 \end{bmatrix} \begin{bmatrix} 1 \\ .5 \\ 1 \\ .25 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ .2 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 1 \\ 1 \\ 0 \end{bmatrix} \end{bmatrix}$$

The Convolution


- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

RGB Image $4 \times 4 \times 3$

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ .5 \end{bmatrix} \begin{bmatrix} .2 \\ .2 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .5 \\ .25 \\ .25 \\ 0 \end{bmatrix} \begin{bmatrix} .15 \\ .1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 1 \\ .25 \\ 1 \\ .25 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ .2 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \\ \begin{bmatrix} 0 \\ .25 \\ .25 \\ 0 \end{bmatrix} \begin{bmatrix} .15 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} .25 \\ .1 \\ .25 \\ .75 \end{bmatrix} \begin{bmatrix} 1 \\ .5 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .25 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \end{bmatrix}$$


The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!



The Convolution



Convolution Hyperparameters

- Number of Filters

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”


Stride 2

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 2



Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 2

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - What is the dimension of the output for Stride 1 vs. Stride 2?

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 5 & 1 \\ 0 & 1 & 2 \\ 1 & 1 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
 - What is output dimension here if stride = 1?

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 5 & 1 \\ 0 & 1 & 2 \\ 1 & 1 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Problem: size of output keep shrinking!
 - Only a few convolutional layers before the resulting 2D dimensions are very small

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Problem: size of output keep shrinking!
 - Only a few convolutional layers before the resulting 2D dimensions are very small
- Solution: Zero padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Padding

Padding
by one

The diagram illustrates a 1D convolution operation. On the left is a 6x6 input matrix with a red dashed box highlighting the top-left 2x2 region (padding of 1). This region is connected by blue dashed lines to a 2x2 output filter matrix on the right, which is highlighted with a blue dashed box. A multiplication symbol (*) is placed between the two matrices. An arrow points from the text 'Padding by one' to the top-right corner of the input matrix, indicating the padding value.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Common choices for a Conv-Layer:
 - Stride = 1
 - Odd Filter Size (3x3, 5x5, etc.)
 - “Same” padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & .5 & 2 \end{bmatrix}$$

Convolution Hyperparameters


- Common choices for a Conv-Layer:
 - Stride = 1
 - Odd Filter Size (3x3, 5x5, etc.)
 - “Same” padding

The diagram illustrates a 1D convolution operation with 'Same' padding. The input is a 6x6 matrix, and the kernel is a 3x3 matrix. The output is also a 6x6 matrix. The top-left 3x3 region of the input is highlighted in red, and the kernel is highlighted in blue. Blue dashed lines show the receptive field of the top-left output element, which is the 3x3 red region.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & .5 & 2 \end{bmatrix}$$


Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun



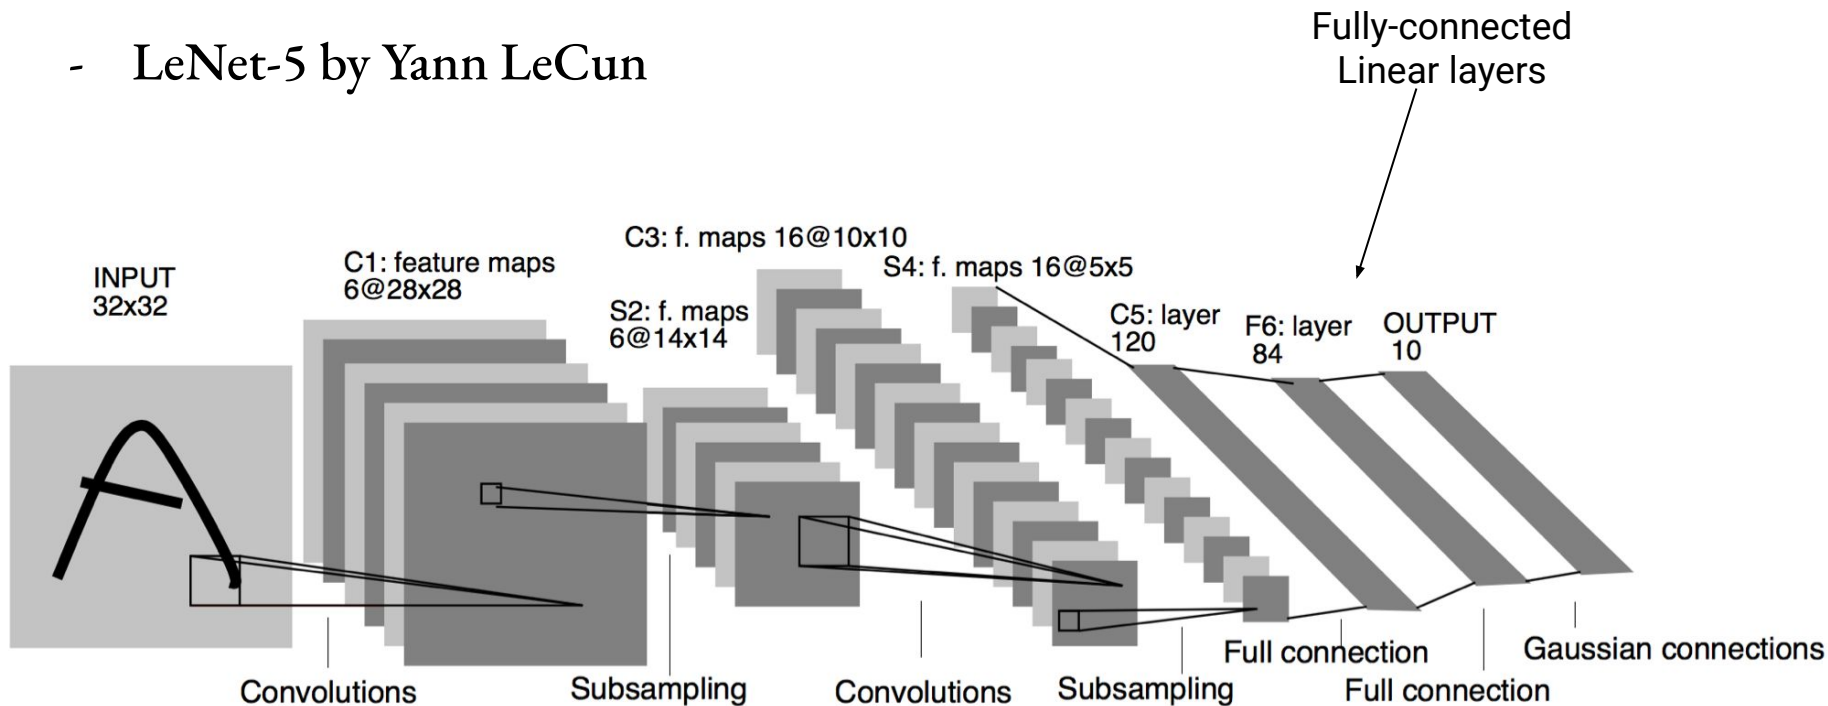
Convolutional Neural Network

- LeNet-5 by Yann LeCun



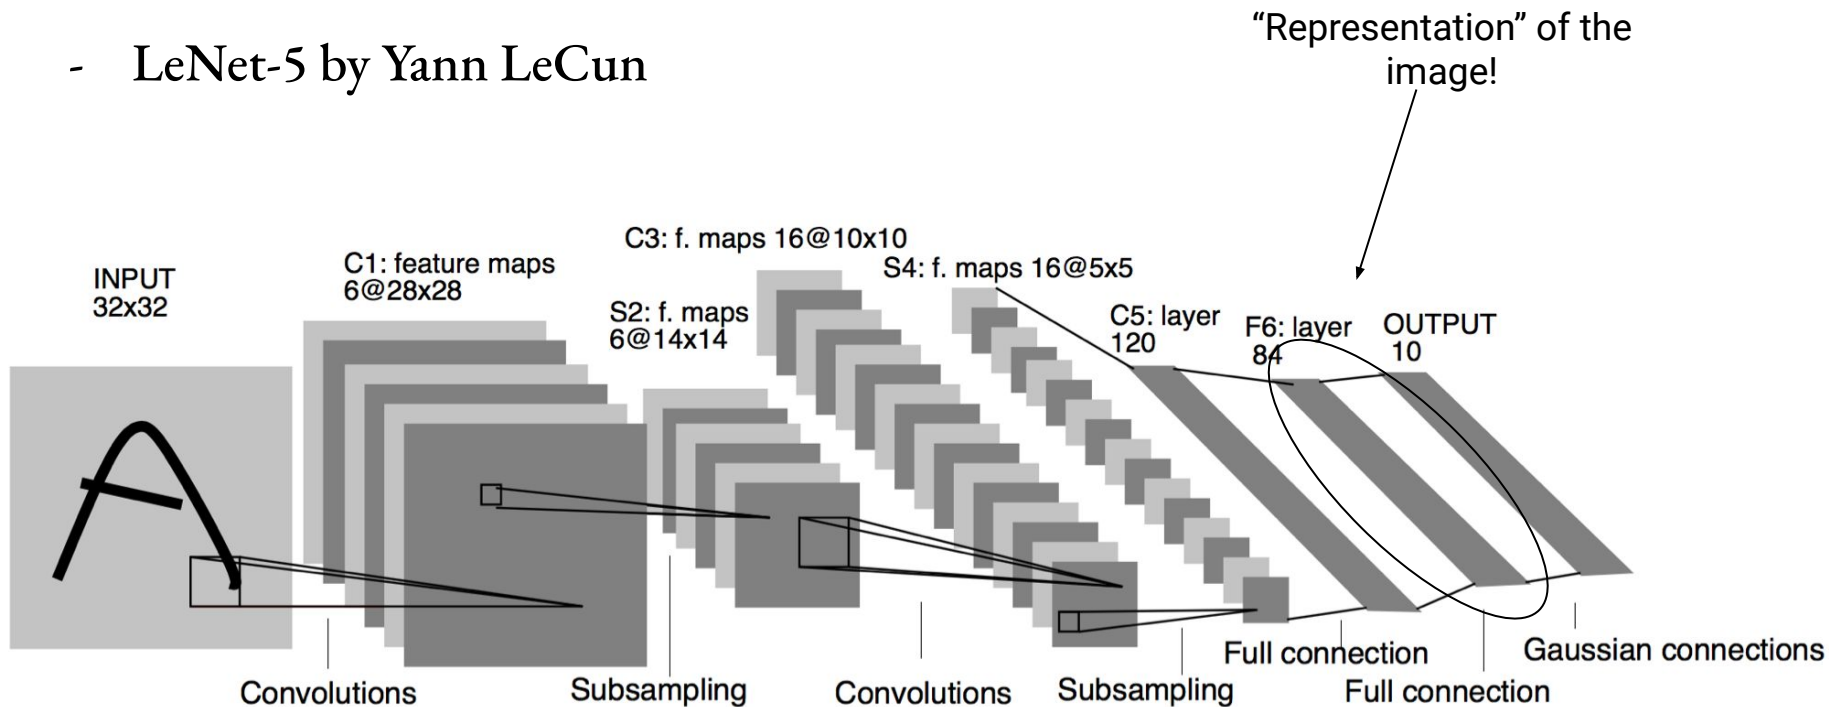
Convolutional Neural Network

- LeNet-5 by Yann LeCun



Convolutional Neural Network

- LeNet-5 by Yann LeCun



Convolutional Neural Network

- AlexNet wins ImageNet Competition in 2012




Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network’s input is 150,528-dimensional, and the number of neurons in the network’s remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Convolutional Neural Network

- AlexNet wins ImageNet Competition in 2012
- By 2015 we have CNNs with >100 layers, better than human-level performance





Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Downsampling

- Reduce size of output
- Minimal information loss in practice
- Intuition: reduce resolution of the image


Downsampling

- Reduce size of output
- Minimal information loss in practice
- Intuition: reduce resolution of the image
- Max Pooling




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
 - Average Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Done along spatial dimension, preserves channels



Convolutional Neural Network

- LeNet-5 by Yann LeCun




Summary

- Convolution Layers
 - Suited for Spatial Data
 - Less Parameters than FC Layers, Weight sharing
- Common Hyperparameters
 - Number of Filters, Filter Size, Stride, Padding
- Common Sequence
 - Conv -> Activation -> Conv -> Activation -> Downsampling
 - Repeat until unrolled into final FC layers


Deeper NNs

- After the success of AlexNet, CNNs got deeper
- Why not just start with as many layers as possible?




Deeper NNs

- After the success of AlexNet, CNNs got deeper
- Why not just start with as many layers as possible?
 - Computer power, data




Deeper NNs

- After the success of AlexNet, CNNs got deeper
- Why not just start with as many layers as possible?
 - Computer power, data
 - Problems with training (vanishing/exploding gradients)



Vanishing/Exploding Gradients




$$a^{[1]} = h(W^{[1]}x + b^{[1]})$$

$$a^{[2]} = h(W^{[2]}a^{[1]} + b^{[2]})$$

$$F(x; \theta) = h(W^{[3]}a^{[2]} + b^{[3]})$$

Vanishing/Exploding Gradients



$$a^{[1]} = h(W^{[1]}x + b^{[1]})$$

$$a^{[2]} = h(W^{[2]}a^{[1]} + b^{[2]})$$

$$F(x; \theta) = h(W^{[3]}a^{[2]} + b^{[3]})$$

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$\frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$\frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$
$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$\frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$

$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$

$$\frac{\partial F}{\partial w_3} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial f_3} \frac{\partial f_3}{\partial w_3}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$(.1)^3 = .001 \quad \frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$
$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$
$$\frac{\partial F}{\partial w_3} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial f_3} \frac{\partial f_3}{\partial w_3}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$(2)^3 = 8 \qquad \frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$

$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$


$$\frac{\partial F}{\partial w_3} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial f_3} \frac{\partial f_3}{\partial w_3}$$

Deeper NNs

- Early parameters can either get stuck, or become unstable during training


Deeper NNs

- Early parameters can either get stuck, or become unstable during training
- Skip Connection




Deeper NNs

- Early parameters can either get stuck, or become unstable during training
- Skip Connection
 - Gradient of earlier parameters depends more directly on output

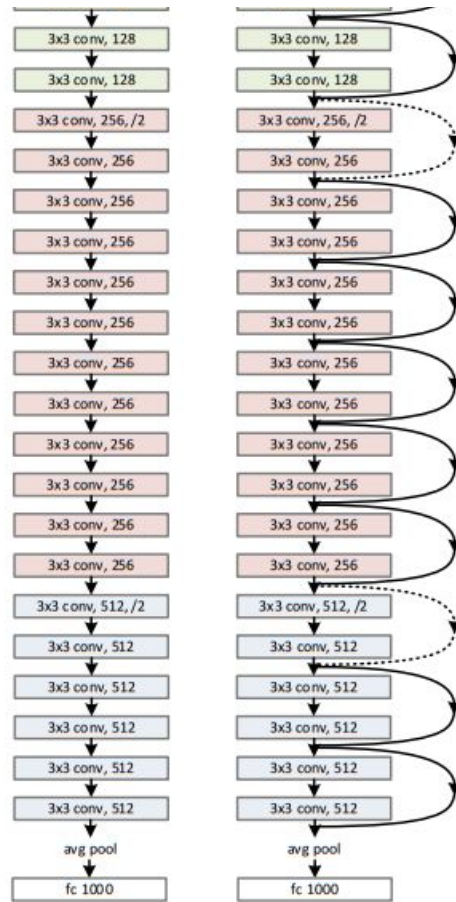






Deeper NNs


- Early parameters can either get stuck, or become unstable during training
- Skip Connection
 - Gradient of earlier parameters depends more directly on output
 - Identity function easier to learn




VGG vs. ResNet



VGG vs. ResNet



“Residual Block”



Other Techniques

- 1x1 Convolutions
 - With a 1x1 filter size you can condense the channel dimension
- Up-convolution
 - “Up-sample” to increase resolution using parameters
 - UNet
- Adaptive Pooling for Fully Convolutional Networks (FCNs)
 - Pool different shaped images to get same size output
- Normalization
 - Batch Normalization, Layer Normalization, Group Normalization
- 1D/3D Convolutions
 - For 3D: filter size maybe 3x3x3, input is of size (C,H,W,L)

UNet

